

# Fortran for Scientific Computing

May 29, 2017

Ge Baolai  
SHARCNET  
Western University

## Outline

- Highlights of the language
- Coding 1 – tryout
- Coffee break
- Coding 2 – 1D diffusion equation
- Parallel computing with coarrays
- Coding Test – Lenna

# SUMMERSCHOOL 2017



```
! Put pieces together

if (this_image() == 1) then
  pic(1:nx,1:ny) = pic_p
  do i = 2, num_images()
    row = (i-1) / g
    col = mod((i-1), g)
    i1 = row*ny_p + 1
    i2 = i1 + ny_p[i] - 1
    j1 = col*nx_p + 1
    j2 = j1 + nx_p[i] - 1
    pic(i1:i2,j1:j2) = pic_p(:, :) [i]
  enddo

  img = trim(base_name)//'.'//'.pgm'

  open(11,file=img,status='unknown')
  write(11,'("P2")')
  write(11,'(i3,i4)' ) nx, ny
  write(11,'(i3)' ) depth
  write(fmt,'(i3)' ) nu-1
  fmt = '(i3,'//trim(fmt)//'i4)'

  do i = 1, ny
    write(11,fmt) pic(i,:)
  enddo
  close(11)
endif
```

You may get the course materials from the URL below:

- To your laptop using browser  
`http://www.sharcnet.ca/~bge/ss2017/fortran_ss2017.zip`
- Or, while on a SHARCNET system, use `wget`  
`wget http://www.sharcnet.ca/~bge/ss2017/fortran_ss2017.zip`
- Create a directory `ss2017` (optional)
- Uncompress the file under Unix environment (in folder `ss2017` if you have one)
- In folder `fortran`, source the `setup.sh` file  
`./setup`

# **Language Basics**



```

DISK OPERATING SYSTEM/360 FORTRAN 360N=FD-451 CL
C ROBERT GLASER, RANDALLSTOWN SENIOR, GROUP A, P AND S
C PRIME NUMBERS
DO 100 I=1,1000
  J=2
  K=2
  2 L=J*K
  IF (L-1) 10,100,10
10 M=2+3
  IF (K-1) 20,3,3
20 K=K+1
  GO TO 2
3 K=2
  IF (J-1) 5,4,4
5 J=J+1
  GO TO 2
4 WRITE (3,6) I
6 FORMAT (I10)
100 CONTINUE
STOP
END

```

# *A brief history...*

## Milestones

- 1954-1957 The birth of FORTRAN by a team led by J. W. Backus at IBM on IBM 704
- FORTRAN II, III, IV and 66
- FORTRAN 77 – a milestone: block structure introduced in the wave of *structured programming*.
- Fortran 90 – a milestone: free format, lots of new features that rejuvenated the language. Fortran more often used instead of FORTRAN. Many Fortran programs we see today were written in the mixed styles of FORTRAN 77 and Fortran 90.
- Fortran 95 – a minor revision, influenced by the *object-oriented programming* concept.
- Fortran 2003 – a minor revision.
- Fortran 2008 – a minor revision: Coarray added for parallel processing.
- Fortran 2015 – to come out in 2016.

## Old FORTRAN (fixed format)

```

1234567 SOURCE STARTS COLUMN 7
      PROGRAM ARRAY
      INTEGER I, J, M, N
      REAL*8 A(100,100)
C
      DO 10 I=1,100
10    B(I)=I
C
C COMMENTS START IN COLUMN 1
      DO 10 J=1,N
      DO 11 I=1,M
      A(I,J)=I*J
11    CONTINUE
10    CONTINUE
C
      DO 20 I=1,100
      A(I,100)=B(I)
20    CONTINUE
      PRINT *, B
      STOP
      END

```

## Modern Format (free format)

```

! Comments start with an !
program array
  integer :: i, j, m, n
  real(8) :: a(100,100), b(100)

  b = [(i,i=1,100)]

  ! Comment can start anywhere
  do concurrent(j=1:n,i=1:n)
    a(i,j) = i * j
  end do

  ! Block assignment like MATLAB
  a(:,100) = b

  print *, b
end program array

```

## Why Fortran

- Designed for scientific computing.
- Simple, takes less time to write (why do you use MATLAB?)
- Performance.
- Widely used and supported.
- Potentially easier for writing parallel code.



- **Integer.**
- **Real** – single or double precision.
- **Complex** – single (8 bytes) and double precision (16 bytes).
- **Character** – character and strings.
- **Logical** – e.g. **.FALSE.** and **.TRUE.**
- **Constant** – literal constants, e.g. 123, 3.1415926...
- **Derived types** – data structures, along with numeric *polymorphism*, lead to **object-oriented** programming
- **Pointer** – useful for array referencing and swapping.

## Examples

**integer**(2) :: n2      ! Short

**integer**(4) :: n      ! Integer

**integer**(8) :: ln      ! Long

**real**(4) :: mass      ! Single precision

**real**(8), **dimension**(10000) :: x, y

**logical** :: isneighbour = .false.

! Coarrays, with [ ], globally accessible

**integer** :: num\_points[\*]

**real, allocatable** :: a(:, :)[:], b(:, :)[:]

**type** particle

real :: m

real :: x, y, z

real :: u, v, w

**end type** particle



- **selected\_real\_kind( $N$ )**, e.g.  
    `real(kind=selected_real_kind(15)) :: v15`  
    `real(kind=selected_real_kind(7)) :: v7`
- **selected\_int\_kind( $N$ )**, e.g. to hold large integers up to 38 digits  
    `integer(kind=selected_int_kind(38)) :: n1, n2, n3`

1

! Non base-10 literal constants

0

b'01100110'      ! Base 2

-9.78654321

o'076543'      ! Octal

+11

z'10fa'      ! Hexadecimal

1.02e-4

-1.0d+3

! Complex constants

$z = (-1, 1)$

$z = (0.123, .99e-2)$

$z = (3.0d0, -4.0d0)$

! Specified precision in N decimal digits

integer, parameter:: **k6**=selected\_int\_kind(6)

-12345\_**k6**

+2\_**k6**

## *Elemental* Mathematical

|                                   |                          |
|-----------------------------------|--------------------------|
| abs(x)                            | !  x                     |
| aimag(z)                          | ! Im(z)                  |
| real(z)                           | ! Re(z)                  |
| int(a), nint(a)                   |                          |
| conjg(z)                          | ! Conjugate              |
| cmplx(x,y)                        | ! x+i*y                  |
| aint(a)                           | ! Nearest w.n. towards 0 |
| anint(a)                          | ! Nearest w.n.           |
| ceiling(a), floor(a)              |                          |
| sqrt(x), exp(x), log(x), log10(x) |                          |
| sin, cos, tan,...                 | ! Trig functions         |
| max/min(a, b[, ...])              |                          |
| mod(a,p), modulo(a,p)             |                          |
| sign(a,b)                         | a  sgn(b)                |
| ... ..                            |                          |

**Unified name for all types**

## Special Functions

|                |   |
|----------------|---|
| erf(x)         |   |
| erfc(x)        |   |
| gamma(x)       |   |
| log_gamma(x)   |   |
| bessel_j0(x)   | ! Bessel, 1 <sup>st</sup> kind, order 0 |
| bessel_j1(x)   | ! Bessel, 1 <sup>st</sup> kind, order 1 |
| bessel_jn(n,x) | ! Bessel, 1 <sup>st</sup> kind, order n |
| bessel_y0(x)   | ! Bessel, 2 <sup>nd</sup> kind, order 0 |
| bessel_y1(x)   | ! Bessel, 2 <sup>nd</sup> kind, order 1 |
| bessel_yn(n,x) | ! Bessel, 2 <sup>nd</sup> kind, order n |
| hypot(x,y)     | $\sqrt{x^2 + y^2}$                      |
| norm2(x)       | $\ x\ _2$                               |
| norm2(x,dim)   |   |

## Inquiry Functions

digits(x)    ! Significant digits in machine rep

epsilon(x)

huge(x)

tiny(x)

maxexpoent(x), minexponent(x)

precision(x)            ! Decimal precision

radix(x)                ! Base

range(x)                ! Decimal exponent

size(a[,dim=1[,...]])    ! Length of an array

shape(a)                ! Dimensions of a

reshape(a,shape)        ! Same as MATLAB

## Array Functions

|                           |                       |
|---------------------------|-----------------------|
| dot_product(x,y)          | $x \cdot y$           |
| matmul(A,B)               | $AB$                  |
| transpose(A)              | $A^T$                 |
| sum(a)                    | $\sum a_i$            |
| product(a)                | $\prod a_i$           |
| maxval(a), minval(a)      |                       |
| maxloc(a), minloc(a)      |                       |
| size(a[,dim=1[,...]])     |                       |
| shape(a)                  | ! Dimensions of a     |
| reshape(a,shape)          | ! Same as MATLAB      |
| cshift(a,p), eoshift(a,p) | ! Shift elements by p |
| pack(dest,mask,src)       |                       |
| unpack(src,mask,dest)     |                       |
| merg(tsrc,fsrc,mask)      |                       |

## Character, String Functions

|                        |   |
|------------------------|---|
| len(s)                 | ! Length of string s  |
| trim(s)                | ! Remove trailing blanks  |
| adjustl(s)             | ! Leading blanks deleted  |
| adjustr(s)             | ! Trailing blanks deleted   |
| repeat(s,n)            | ! Repeat s n times  |
| index(s, t[, back])    | ! Starting position of t  |
| scan(s, set[, back])   | ! Position of a char in set   |
| verify(s, set[, back]) | ! 0 if each char in s is in set; or the position of a char not in set |
| achar(i)               | ! ASCII char  |
| iachar(c)              | ! Position in ASCII   |

*No one can remember all of them. Always keep a reference handy.*

## Time Functions

! Equivalent to `gettimeofday()`

```
call date_and_time(  
    date,  
    time,  
    zone,  
    value  
)
```

! Measures elapsed time, mostly used

```
call system_clock(  
    count,  
    count_rate,  
    count_max  
)
```

! Less used

```
call cpu_time(time)
```

## Random Numbers

! Seeds are in an array, get its size

```
call random_seed(size = lseed)
```

! Create an array for seeds

```
allocate(seed(lseed))
```

! Create seeds with clock ticks

```
call system_clock(count=clock)
```

```
seed = clock + 37 * [(i - 1, i = 1, lseed)]
```

! Plant the seeds

```
call random_seed(put=seed)
```

! Generate a sequence of RN's

```
call random_number(r)
```

## Scalars

$x + y$

$x - y$

$x * y$

$a/b$

$x^{**}y$       !  $x^y$

$n = 13.4$       !  $n = 13$

$i = 13/2$       !  $i = 6$

$\text{real} :: a, b$

$\text{complex} :: z$

$z = \text{cmplx}(a, b)$       ! convert to complex

$z = (-3.3, 4.6)$       ! or do this way

$a = z$       ! a holds  $\text{Re}(z)$

$c = (-3.3, 4)$       ! c holds -3.3

## Structure

**type interval**      ! Define a new type

$\text{real} :: \text{lo}$

$\text{real} :: \text{up}$

**end type interval**

**type(interval) :: w, u, v**

$\text{real} :: a, b$

$w = \text{interval}(-1.0, 1.0)$  ! Initialize a structure

$a = w \% \text{lo}$

$b = w \% \text{up}$

$u = w$       ! Structure assignment

$v = 2*w - u$       ! Not yet defined



## Scalar, Array Operations

```
pi = 3.1415926
```

```
c = 2.0*pi*r
```

```
a = b * c + d / e
```

```
r = sqrt(x**2 + y**2)
```

```
integer :: m = 17, n = 5
```

```
real :: q, a
```

```
q = m / n
```

```
a = q * n
```

**! What's in a?**

## Character, Strings

```
character(len=5) :: word1
```

```
character(len=5) :: word2
```

```
character(256) :: grtg_msg
```

## String concatenation

```
word1 = 'Hello'
```

```
word2 = 'world'
```

```
grtg_msg = word1//', '//word2//'!'
```

## Trimming off trailing blanks

```
trim(grtg_msg)
```

## Objects of Derived Type

```
type particle
```

```
  real :: m
```

```
  real :: x, y
```

```
  real :: u, v
```

```
end type particle
```

```
type(particle) :: p, q
```

```
p = particle(0.2, -1.0, 3.5, 0.5, 2.7)
```

```
q = p
```

```
q%x = q%x - x0
```

```
q%y = q%y - y0
```

## Overloading Operators

! Define an interface

```
interface operator (+)
```

```
  function rational_add(x,y)
```

```
    type(rational) :: rational_add
```

```
    type(rational) :: x, y
```

```
  end function rational_add
```

```
end interface
```

! Define the function

```
function rational_add(x,y)
```

```
  type(rational) :: rational_add
```

```
  type(rational) :: x, y
```

```
  ... ..
```

```
end function rational_add
```

## Arrays

```
A = [1, 2, 3, -1, 4.5]
```

```
A = B
```

### ! Traditional loop operations

```
do j = 1, 1000
```

```
  do i = 1, 1000
```

```
    a(i,j) = b(i,j) + c(i,j)
```

```
  end do
```

```
end do
```

### ! Whole array operations

```
a = b + c
```

```
y(1:n) = x(1:n)
```

```
c(1,:) = a(:,1)
```

## Pointers

```
real, target, allocatable :: ws_u(:, :), ws_v(:, :)
```

```
real, pointer, dimension(:) :: u, v, tmp
```

### ! Allocate spaces for arrays

```
allocate(ws_u(n,n),ws_v(n,n))
```

### ! Make pointers point to arrays

```
u => ws_u; v => ws_v
```

### ! Use pointers as if arrays

```
uxx = (u(i-1,j) - 2.0*u(i,j) + u(i+1,j))/(dx*dx)
```

```
uyy = (u(i,j-1) - 2.0*u(i,j) + u(i,j+1))/(dy*dy)
```

### ! Swap pointers instead of arrays

```
tmp => u; u => v; v => tmp
```

## **DO loops**

```
do i = 1, n  
    y(i) = f(x(i))  
end do  
  
do k = 1, 10000, 2  
    do something  
end do  
  
do j = 1, n  
    do i = 1, m  
        a(i,j) = mat_setval(i,j)  
    end do  
end do
```

## **DO [ WHILE ] loops**

```
eps = 1.0  
do while (eps + 1.0 /= 1.0)  
    eps = eps / 2  
end do  
  
do  
    do something  
    if (abs(error) <= tol) exit  
end do  
  
do  
    read *, x  
    if (x < 0) cycle  
    exit  
end do
```

## **FORALL**

```
forall (i=1:m, j=1:n, y(i,j) /= 0.)  
    a(j,i) = 1.0/y(i,j)  
end forall
```

## **DO CONCURRENT**

```
do concurrent (i=1:m, j=1:n)  
    a(i,j) = a(i,j) + alpha*b(i,j)  
enddo  
  
do concurrent (i=1:m, j=1:n, i/=j)  
    ... ..  
enddo
```

## **Coarray**

- A parallel programming paradigm without explicit interprocess communication calls, e.g. via MPI.
- Data objects globally accessible amongst processes without explicit data transfer operations, e.g. MPI\_Send(), MPI\_Recv(), etc.
- Uses one-sided communication model, with “fetch” and “push” to get and put data respectively.
- Supports parallel processing on both shared and distributed memory architectures.
- In favour of thinking of algorithms than implementation tediousness.

## *IF condition statement*

The relational operators

.lt.    or    <    less than

.gt.    or    >    greater than

.le.    or    <=    less than or equal

.ge.    or    >=    greater than or equal

.eq.    or    ==    equal

.ne.    or    /=    not equal

Logical expressions

.not.    Negation

.and.    Logical intersection

.or.    Logical union

## *IF..THEN..ELSE IF..ELSE..ENDIF*

**if** (x == y) **then**

*do something*

**end if**

**! Also allow arrays**

**if** (x == y1) **then**

*do case1*

**else if** (x == y2) **then**

*do case2*

**else if** (x == y3) **then**

*do case3*

**else**

*default action*

**end if**

# **Coding Exercise 1 – *Tryout***



## **SELECT..CASE**

```
select case expr  
case val1  
    process case1  
case val2  
    process case2  
...  
case default  
    process default case  
end select
```

## **Example**

**! Select on individual values**

```
select case j  
case 1  
    call sub1  
case 2  
    call sub2  
end select
```

**! Select on a range**

```
select case x  
case (:-1) ! All <= -1  
    call sub1  
case (1:5) ! 1,2,3,4,5  
    call sub2  
...  
end select
```

## **ANY, ALL**

```
if (any(a > b)) then
    print *, 'An elem in A > an elem in B'
endif
```

```
if (all(a > b)) then
    print *, 'All elem in A > elems in B'
else if (all(b > a)) then
    print *, 'All elem in B > elems in A'
else
    print *, 'Set A and B overlap'
endif
```

## **WHERE [ ..ELSEWHERE..ENDWHERE ]**

```
where (b < 0) a = b
print *, 'where (b<0) a=b:'
print '(10f5.0)', (a(i), i = 1,10)
```

```
where (b /= 0)
    a = a/b
elsewhere
    a = 0
endwhere
```

## **Mask**

### ! Using loop

```
do concurrent(i = 1:n, a(i) > b(i))
```

```
    a(i) = a(i) - b(i)*d(i)
```

```
    c(i) = c(i) + a(i)
```

```
enddo
```

### ! Using logical mask and merge()

```
logical:: mask(n)
```

```
... ..
```

```
mask = a > b
```

```
a = a - merge(b*d,0.,mask)
```

```
c = c + merge(a,0.,mask)
```

- If branches can be avoid.
- Some compilers might be able to better job.

## MATLAB

```
v = [-1, 0, 2 5, 14, 3.5, -0.02];
```

```
n = 1000;
```

```
x = 1:n;
```

```
y = cos(x)
```

## Fortran

```
v = [-1, 0, 2 5, 14, 3.5, -0.02]
```

```
v = (/ -1, 0, 2, 5, 14, 3.5, -0.02 /)
```

```
n = 1000
```

```
x = [(i,i=1,n)]
```

! or `x = (/ (i,i=1,n) /)`

```
y = cos(x)
```

! `x, y` are arrays

## MATLAB

```
v = 1:n;  
a = reshape(v, n1, n2);
```

% It's in column order

|   |   |    |    |
|---|---|----|----|
| 1 | 5 | 9  | 13 |
| 2 | 6 | 10 | 14 |
| 3 | 7 | 11 | 15 |
| 4 | 8 | 12 | 16 |

## Fortran

```
v = [i=1,n]  
a = reshape(v, [n1,n2])
```

! It's in row order

|    |    |    |    |
|----|----|----|----|
| 1  | 2  | 3  | 4  |
| 5  | 6  | 7  | 8  |
| 9  | 10 | 11 | 12 |
| 13 | 14 | 15 | 16 |

## MATLAB

```
asize = size(a);  
n1 = asize(1);  
n2 = asize(2);
```

## Fortran

```
asize = shape(a)  
n1 = asize(1)  
n2 = asize(2)  
  
nrows = size(a,dim=1)  
ncols = size(a,dim=2)  
num_elems = size(a)  
  
x_len = size(x)
```

## MATLAB

```
a = zeros(100,100);  
b = zeros(100);
```

```
n = input("Enter n (n <= 100): ");
```

```
a(:, :) = 1.0;  
b(:) = 2.0;  
c = a(1,:) ./ b;
```

```
a(::2) % Get even index elements.  
a(5:1:-1) % Traverse in reverse order.
```

## Fortran

```
real, dimension(100,100) :: a  
real, dimension(100) :: b, c
```

```
print *, 'Enter n (n <= 100):'  
read *, n
```

```
a = 1.0  
b = 2.0  
c = a(1,:) / b
```

```
a(::2) ! Get event index elements.  
a(5:1:-1) ! Traverse in reverse order.
```



## MATLAB

```
a = [1,2,3,4,5,6,7,8,9,10]
```

```
a(3)
```

```
% N/A
```

```
% N/A
```

```
a(2:5)
```

```
% N/A
```

```
% N/A
```

```
% N/A
```

```
a(:)
```

## Fortran

```
a = [1,2,3,4,5,6,7,8,9,10]
```

```
print '(10i5)', a(3)    ! 3
```

```
print '(10i5)', a(4:)   ! 4 5 6 7 8 9 10
```

```
print '(10i5)', a(:6)   ! 1 2 3 4 5 6
```

```
print '(10i5)', a(2:5)  ! 2 3 4 5
```

```
print '(10i5)', a(::3)  ! 1 4 7 10
```

```
print '(10i5)', a(2:4:2) ! 2 4
```

```
print '(10i5)', a(1::2) ! 1 3 5 7 9
```

```
print '(10i5)', a(:)    ! 1 2 3 4 5 6 7 8 9 10
```

## Format

```
a([istart]:[iend]][:incr])
```

## New functions

- `dot_product(a, b)`  
 $\langle a, b \rangle$
- `product(a)`  
$$\prod_{i=1}^n a_i$$
- `sum(a)`  
$$\sum_{i=1}^n a_i$$
- `maxcal(a)/minval(a)`

## Common intrinsic apply to arrays

- SIN, COS, etc apply to array arguments as well.

## **Coding Exercise 2 – *Diffusion Equation***

Solve

$$\frac{\partial u}{\partial t} = a \frac{\partial^2 u}{\partial x^2}$$

with initial and boundary conditions

$$\begin{aligned} u(x, 0) &= f(x), \\ u(-L, t) &= u(L, t) = 0. \end{aligned}$$

**Numerical solution:**

Using a grid with time step  $\Delta t$  chosen properly and  $\Delta x = 2L/(N - 1)$



$$\begin{aligned} x_i &= \Delta x(i - 1), \quad i = 1, \dots, N, \\ t_n &= \Delta t n, \quad n = 1, \dots \end{aligned}$$

Denote the approximation

$$U_i^n \approx u(x_i, t_n)$$

**Exercise: Set**

$$f(x) = \begin{cases} H - \frac{|x|}{\epsilon} & -\epsilon \leq x \leq \epsilon, \\ 0 & \text{otherwise.} \end{cases}$$



- Apply (explicit scheme) finite difference approximation to obtain

$$\frac{U_i^{n+1} - U_i^n}{\Delta t} = a \left( \frac{U_{i-1}^n - 2U_i^n + U_{i+1}^n}{\Delta x^2} \right)$$


- Or, in an iterative form, for  $i = 2, \dots, N - 1$ ,  $n = 0, \dots$

$$U_i^{n+1} = \left( 1 - 2\frac{a\Delta t}{\Delta x^2} \right) U_i^n + \frac{a\Delta t}{\Delta x^2} U_{i-1}^n + \frac{a\Delta t}{\Delta x^2} U_{i+1}^n$$

or simply noted

$$U_i^{n+1} = (1 - 2\lambda)U_i^n + \lambda U_{i-1}^n + \lambda U_{i+1}^n$$


- Compute all  $U_i^{n+1}$  for  $i = 1, \dots, N$ ,  $n = 0, \dots$
- Generate graph of  $U_i^{n+1}$  every 50 or so time steps.



**Homework:** Modify the 1d diffusion code for parallel processing.

Hint: Divide the interval  $[0, L]$  into  $p$  sub-intervals. Assign each to a process (image). Each works on its own, image 1 collects final results. See the slides on coarray for implementation details.

*Serial:* Sweep through all points



*Parallel:* Each image sweeps through all points on its own portion. End points need values from its immediate neighbours.





Image 1

Image 2

Image 3

Image 4


# Parallel Computing with Coarrays





The idea...

- I do not have the portion of A that k has
- I want to copy it from k



Use message passing, we would write

- On rank 1, to receive data from rank k

```
MPI_Recv(A(is:ie,js:je),n,MPI_REAL,k,tag,MPI_COMM_WORLD,status)
```

Or, more generic

```
MPI_Recv(buffer,n,MPI_REAL,k,tag,MPI_COMM_WORLD,status)
```

*Put buffered data into A*

- On rank k, to send data to rank 1


```
MPI_Send(A(is:ie,js:je),n,MPI_REAL,1,tag,MPI_COMM_WORLD)
```

Or

*Copy data from local A to the buffer*

```
MPI_Send(buffer,n,MPI_REAL,1,tag,MPI_COMM_WORLD)
```

- One must ensure the assembly is correct!



k

1

But what we really want is essentially as simple as this...

$$A(is:ie,js:je) \leftarrow A(is:ie,js:je)^k$$

So here come this notion

$$A(is:ie,js:je) = A(is:ie,js:je)[k]$$

program main

real :: x(10000), u(10000)

real :: A(1000,1000)[\*]

complex :: y(10000)

... ..

$$A(i1:i2,j1:j2) = A(i3:i4,j3:j4)[k]$$

end program main

program main

real :: x(10000), u(10000)

real :: A(1000,1000)[\*]

complex :: y(10000)

... ..

$$A(i1:i2,j1:j2)[k] = A(i3:i4,j3:j4)$$

end program main

program main

real :: x(10000), u(10000)

real :: A(1000,1000)[\*]

complex :: y(10000)

... ..

$$A(i1:i2,j1:j2) = A(i3:i4,j3:j4)[k]$$

end program main

## Coarray Syntax

- Globally addressible arrays amongst processes – **images**.
- Each image holds the same size copies of data objects – **coarrays**.
- Data objects with subscripts in square brackets indicates coarray, in any of the following forms
  - $X[*]$       ! Upper bound not set
  - $X[16]$       ! Max images 16
  - $X[p,q]$       ! p-by-q images
  - $X[p,*]$       ! Last bound not set
  - $X[8,0:7,1:*$  ! Three codimensions
- [*identifier*] defines the number of images (and topology)
- Upper bound usually not defined.

## Example

! Array coarrays

```
real :: a(1000,1000)[*]
```

```
real :: b(1000,1000)[16,16], x(10000)[16]
```

```
complex, allocatable, codimension[*] :: z(:)
```

! Scalar coarrays

```
integer :: m[*], n[*]
```

```
if (this_image() == 1) then
```

```
  input data
```

```
  do image = 1, num_images()
```

```
    u[image] = u ! Send u to all images
```

```
  enddo
```

```
endif
```



## Coarray Syntax (cont'd)

- Objects of derived types

`type(type1) :: p[*]`

`type(type2), allocatable :: u[:]`

## Example

**! Derived data types**

**type** particle

`real :: m`

`real :: x, y, z`

`real :: u, v, w`

**end type** particle

**! Static storage**

`type(particle):: p(1000000)[*]`

**! Dynamic storage**

`type(particle), allocatable:: p(:)[:]`

`u = p(k)[16]%u`

`v = p(k)[16]%v`

## Concept

## Images

a=1, b=2

a=2, b=4

a=3, b=6

.  
.  
.

a=16, b=32

## Execution of code

```
do i = 1, num_images()
  print *, a[i], b[i]
enddo
```

## Example

```
program try_coarray
```

```
real :: a[*]      ! Declare a as coarray obj
```

```
real, codimension[*] :: b ! Or this way
```

```
! a and b below are local to the image
```

```
a = this_image()
```

```
b = this_image()*2
```

```
! Access a and b on other images
```

```
if (this_image() == 1) then
```

```
  do image = 1, num_images()
```

```
    print *, 'Image', this_image(), a[i], b[i]
```

```
  enddo
```

```
endif
```

```
end program try_coarray
```

- Access coarray objects by referencing to the object with an image index in square [ ], e.g.

$x[i] = y$                       ! Push local value  $y$  to  $x$  on image  $i$

$a(:, :)[i] = b$                 ! Whole array assignment not used in coarrays

$z = z[i]$                       ! Fetch value of  $z$  on image  $i$  and assign it to local  $z$

- Note the following is executed by every image (due to SPMD model)

$x[16] = 1$

- For selective execution

if ( $this\_image() == 16$ ) then

$x = 1$

endif

- Note Fortran arrays use ( ) for array elements, not [ ], so there is no confusion!

program ex1

implicit none

real :: z[\*]

integer :: i

print '("Image",i4,": before: z=",f10.5)', this\_image(), z

sync all

if (this\_image() == 1) then

read \*, z

do i = 2, num\_images()

z[i] = z

enddo

endif

sync all

print '("Image",i4,": after: z=",f10.5)', this\_image(), z

end program ex1

program ex2

character(80) :: **host**[\*] ! Note: host – local; host[i] – on image i

integer :: i

call get\_environment\_variable("HOSTNAME",value=**host**)

if (this\_image() == 1) then

do i = 1, num\_images()

print \*, 'Hello from image', i, 'on host ', trim(**host**[i])

enddo


endif

end program ex2

- Any comments on the broadcast operation?



```
do i = 2, num_images()  
  z[i] = z  
enddo
```

## Linear



$$O(N)$$

## Improved





- The SPMD model is assumed, i.e. every image executes the same program.
- The SPMD model assumes coarrays on every image, e.g.  
    `real :: a(10000,10000)[*]`  
    `integer :: ma[*], na[*]`
- The SPMD model requires self identification (“this image”) and others, via
  - `this_image()`
  - `num_images()`
- The control of work flow is done by the selection logics, e.g.  
    `if (1 == this_image()) then`  
        `call manager()`  
    `else`  
        `call worker()`  
    `endif`
- Memory coherence is not assured until you want to (e.g. via remote copies)
- Synchronizations

# Coding Test – *Lenna*

## The problem

- Each process (image) reads and posses a small (square) portion of Lenna, labelled sequentially.
- To have the main process collect portions of Lenna and assemble them into the whole image.
- The main process then writes it out to a PGM file.

## The implementation

- Use `pic(:, :)` for the whole and `pic_p(:, :)[ ]` for local portion.
- The main process loops over processes, collects the portion from each process and assembles it in the whole array accordingly.



Source: <https://en.wikipedia.org/wiki/Lenna>

```

if (this_image() == 1) then
  do i = 1, num_images()
    ... ..
    pic(i1:i2,j1:j2) = pic_p[i] ! Fetch the portion from image i
  enddo
endif


```

! Import the portion, e.g.

```
allocate(pic_p(nx_p,ny_p)[*])
do i = 1, ny_p
  read(10,*) pic_p(i,:)
enddo
```

! Perform some tasks

... ..



```
if (this_image() == 1) then
  do i = 1, num_images()
```

... ..

```
pic(i1:i2,j1:j2) = pic_p[i] ! Fetch the portion from image i
```

```
enddo
```

```
endif
```



## The Procedure

- Use the base source provided and complete it;
- Each process – *image* – reads a file containing the distinct portion of Lenna into an array **pic\_p** declared as coarray. This is already implemented.
- Image 1 will fetches a portion of Lenna from each of the rest images and assembles the pieces into whole array **pic**.
- You need to calculate the start and end indices of the rows and columns in **pic** where the fetched picture is to be inserted.
- Image 1 then writes the restored picture stored in **pic** to a PGM file (implemented);
- Compile the program using the Makefile provided to generate executable **lenna**  
**make**
- Run the program with command  
**mpirun -n 4 ./lenna**
- Use Unix command **display** to view your result  
**display Lenna.pgm**

- [1] Michael Metcalf, John Reid and Malcolm Cohen, “***Modern Fortran Explained***”, Oxford University Press, New York, 2011.
- [2] Sun Microsystems, Inc., “***Fortran Programming Guide***”, 2005.
- [3] **JTC1/SC22** – The international standardization subcommittee for programming languages (<http://www.open-std.org/jtc1/sc22/>).