

# MPI Day-2:

Performance/Scaling

Non-blocking communication

Communicators/Groups/Topology

One-side communication

FEI MAO

SHARCNET


# **Performance/Scaling**

# Scaling — Throughput

- ▶ How a problem's throughput scales as processor number increases (“strong scaling”).
- ▶ In this case, linear scaling:

$$H \propto P$$

- ▶ This is **Perfect scaling**.




# Scaling – Time

- ▶ How a problem's timing scales as processor number increases.
- ▶ Measured by the time to do one unit. In this case, inverse linear scaling:

$$T \propto 1/P$$

- ▶ Again this is the ideal case, or “embarrassingly parallel”.




# Scaling – Time

- ▶ How a problem's timing scales as processor number increases.
- ▶ Measured by the time to do one unit. In this case, inverse linear scaling:

$$T \propto 1/P$$

- ▶ Again this is the ideal case, or “embarrassingly parallel”.




# Scaling – Speedup

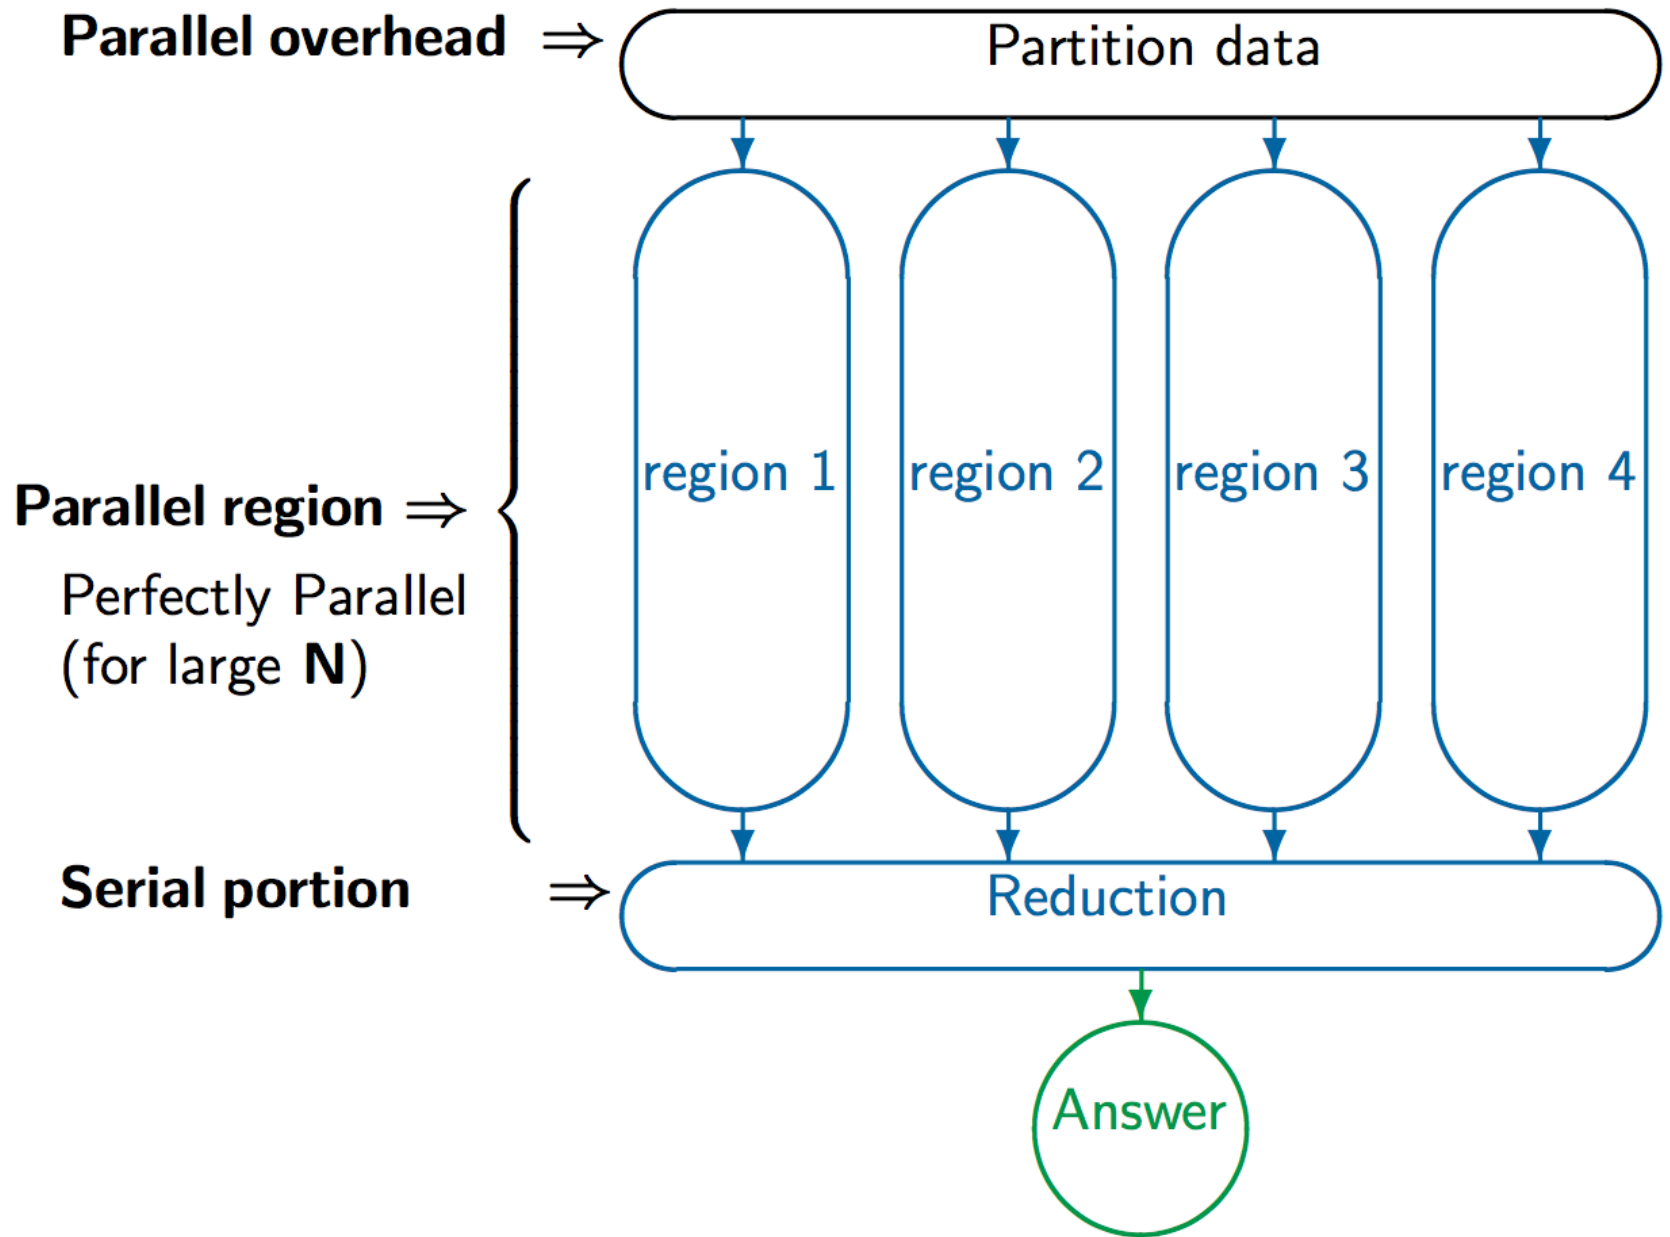
- ▶ How much faster the problem is solved as processor number increases.
- ▶ Measured by the serial time divided by the parallel time

$$S = \frac{T_{\text{serial}}}{T(P)} \propto P$$

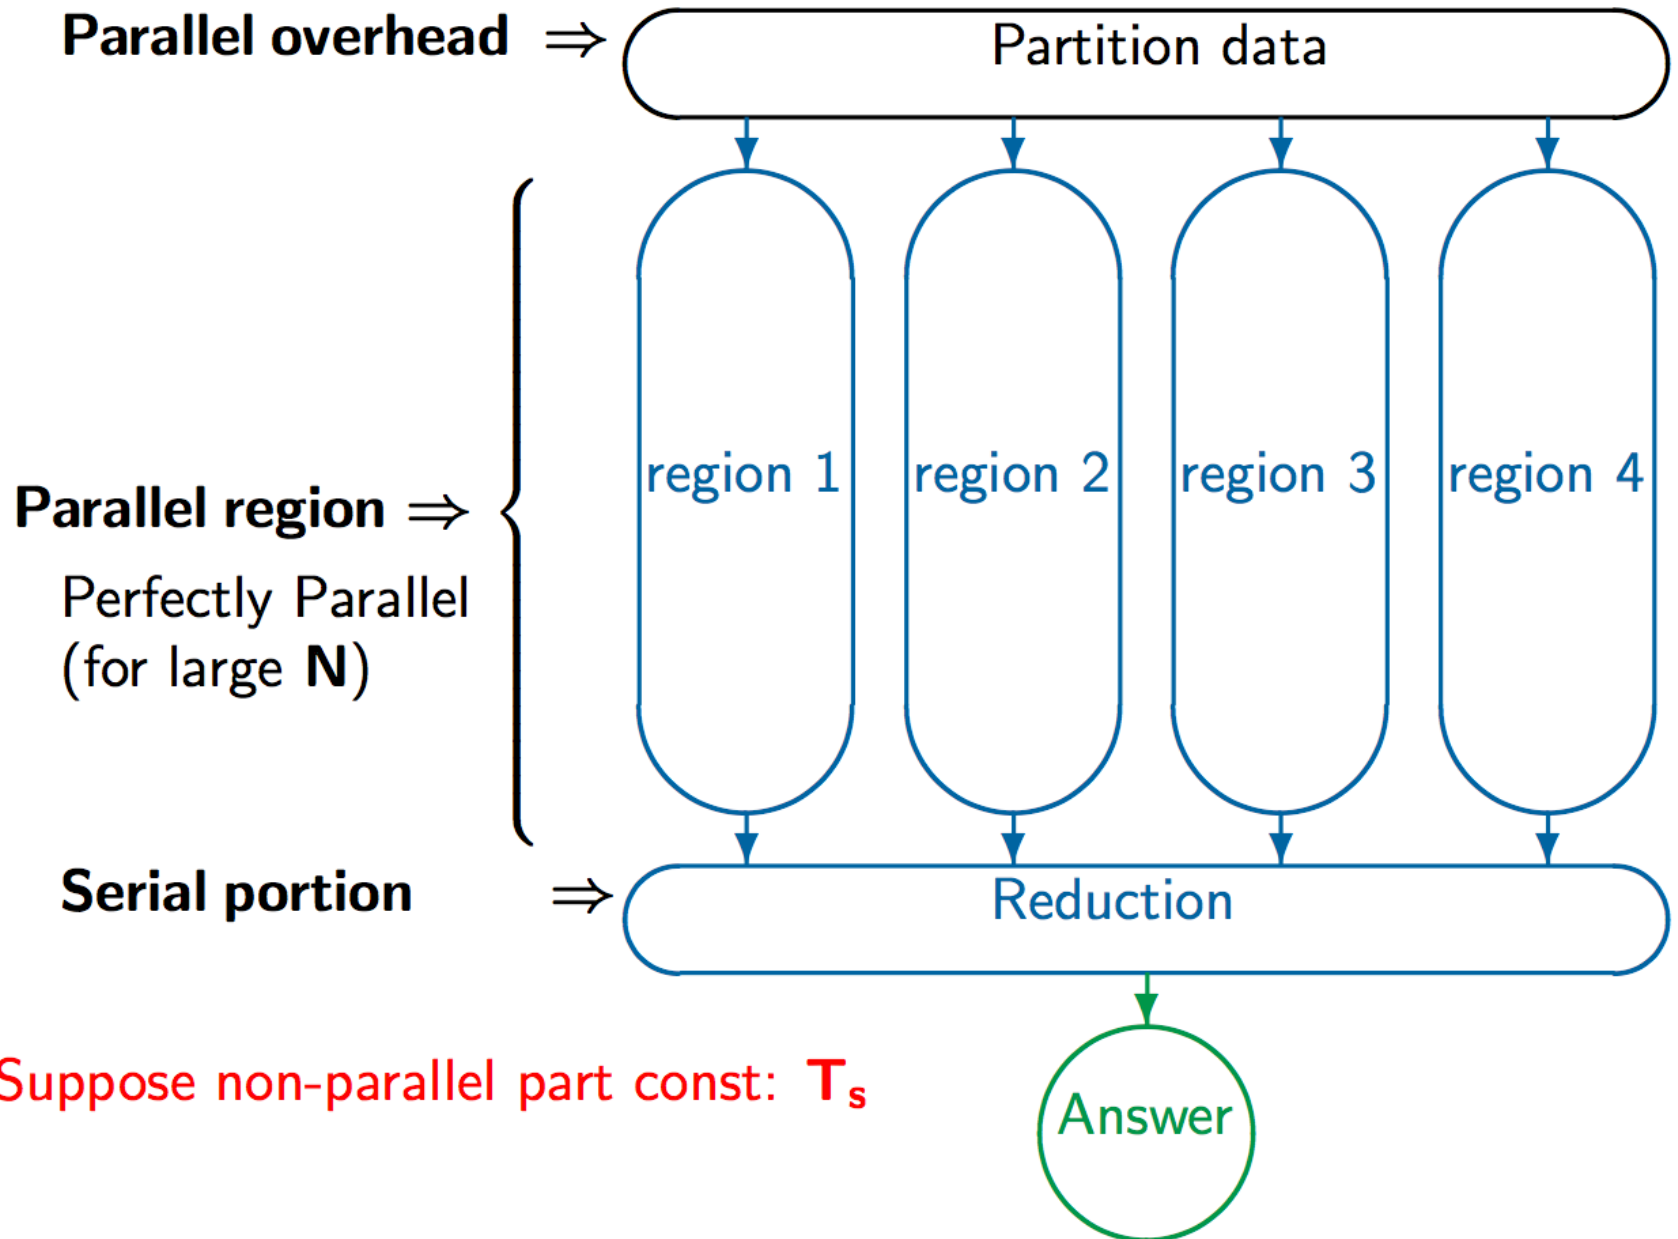
- ▶ For embarrassingly parallel applications: Linear speed up.



# Serial Overhead



# Serial Overhead






# Amdahl's law

Speed-up (without parallel overhead):

$$S = \frac{NT_1 + T_s}{\frac{NT_1}{P} + T_s}$$

or, calling  $f = T_s/(T_s + NT_1)$  the serial fraction,

$$S = \frac{1}{f + (1 - f)/P}$$




# Amdahl's law

Speed-up (without parallel overhead):

$$S = \frac{NT_1 + T_s}{\frac{NT_1}{P} + T_s}$$

or, calling  $f = T_s/(T_s + NT_1)$  the serial fraction,

$$S = \frac{1}{f + (1 - f)/P} \quad \xrightarrow{P \rightarrow \infty} \frac{1}{f}$$



Serial part dominates asymptotically.

Speed-up limited, no matter size of **P**.

And this is the overly optimistic case!


(for  $f = 5\%$ )

# Trying to beat Amdahl's law

## Scale up!

The larger **N**, the smaller the serial fraction:

$$f(P) = \frac{P}{N}$$

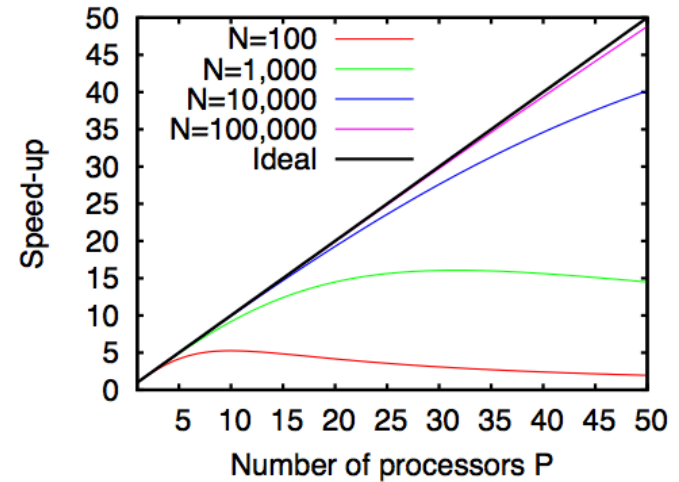


# Trying to beat Amdahl's law

## Scale up!

The larger **N**, the smaller the serial fraction:

$$f(P) = \frac{P}{N}$$



Weak scaling: Increase problem size while increasing **P**

$$\text{Time}_{\text{weak}}(P) = \text{Time}(N = n \times P, P)$$

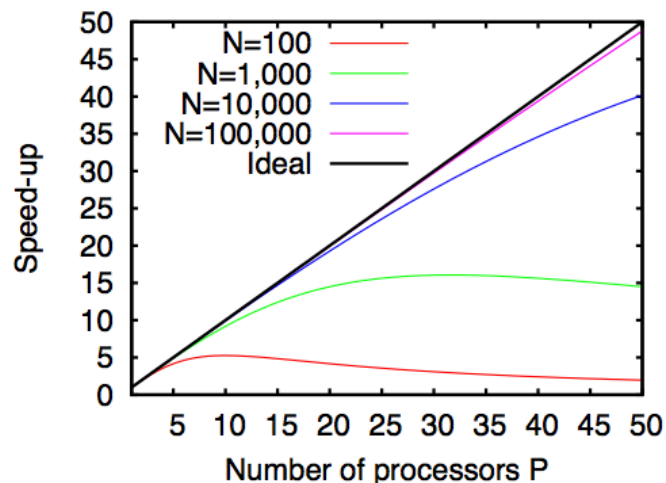
Good weak scaling means this time approaches a constant for large **P**.

# Trying to beat Amdahl's law

## Scale up!

The larger **N**, the smaller the serial fraction:

$$f(P) = \frac{P}{N}$$



Weak scaling: Increase problem size while increasing **P**

$$\text{Time}_{\text{weak}}(P) = \text{Time}(N = n \times P, P)$$


Good weak scaling means this time approaches a constant for large **P**.

## Gustafson's Law

Any large enough problem can be efficiently parallelized (Efficiency  $\rightarrow 1$ ).


# Synchronization Overhead

- ▶ Most problems are not purely concurrent.
- ▶ Some level of synchronization or exchange of information is needed between tasks.
- ▶ While synchronizing, nothing else happens: increases Amdahl's  $f$ .
- ▶ And synchronizations are themselves costly.



# Load Balancing

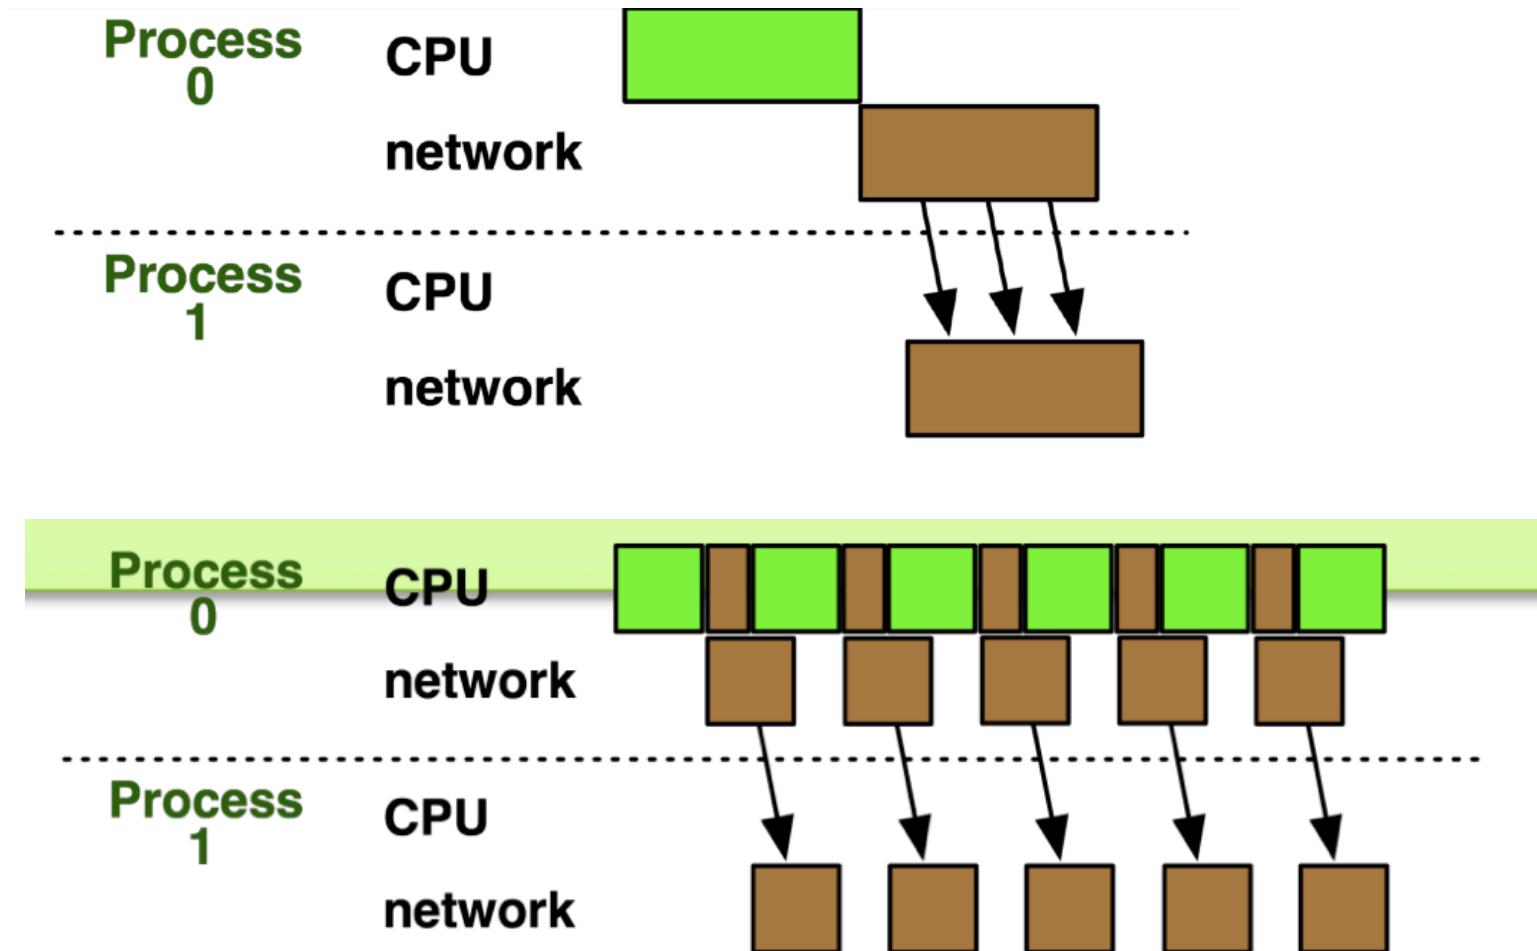
- ▶ The division of calculations among the processors may not be equal.
- ▶ Some processors would already be done, while others are still going.
- ▶ Effectively using less than  $P$  processors: This reduces the efficiency.
- ▶ Aim for load balanced algorithms.



# **NONBLOCKING COMMUNICATION**



# NONBLOCKING COMMUNICATION



# NON-BLOCKING COMMUNICATION


- Non-blocking sends and receives
  - `MPI_Isend` & `MPI_Irecv`
  - returns immediately and sends/receives in background
- Enables some computing concurrently with communication
- Avoids many common dead-lock situations
- Also non-blocking collective operations in MPI 3.0

# NON-BLOCKING COMMUNICATION

- Have to finalize send/receive operations
  - MPI\_Wait, MPI\_Waitall, ...
    - Waits for the communication started with MPI\_Isend or MPI\_Irecv to finish (blocking)
  - MPI\_Test, ...
    - Tests if the communication has finished (non-blocking)
- You can mix non-blocking and blocking p2p routines
  - e.g., receive MPI\_Isend with MPI\_Recv

# Typical usage pattern

- `MPI_Irecv(ghost_data)`
- `MPI_Isend(border_data)`
- `Compute(ghost_independent_data)`
- `MPI_Waitall` `Compute(border_data)`



# Non-blocking send

`MPI_Isend(buf, count, datatype, dest, tag, comm, request)`

- Parameters
  - Similar to `MPI_Send` but has an additional request parameter
- `buf` send buffer that must not be written to until one has checked that the operation is over
- `request` a handle that is used when checking if the operation has finished (integer in Fortran, `MPI_Request` in C)

# Non-blocking receive

`MPI_Irecv(buf, count, datatype, source, tag, comm, request)`

- parameters similar to `MPI_Recv` but has no status parameter
- `buf` receive buffer guaranteed to contain the data only after one has checked that the operation is over
- `request` a handle that is used when checking if the operation has finished

# Wait for non-blocking operation

`MPI_Wait(request, status)`

- Parameters
  - `request` handle of the non-blocking communication
  - `status` status of the completed communication
- A call to `MPI_WAIT` returns when the operation identified by `request` is complete

# Wait for non-blocking operation

MPI\_Waitall(count, requests, status)

- Parameters
  - count number of requests
  - requests array of requests
  - status array of statuses for the operations that are waited for
- A call to MPI\_Waitall returns when *all* operations identified by the array of requests are com



# Additional completion operations

other useful routines:

- **MPI\_Waitany**
- **MPI\_Waitsome**
- **MPI\_Test**
- MPI\_Testall
- MPI\_Testany
- MPI\_Testsome
- MPI\_Probe

# Wait for non-blocking operations

MPI\_Waitany(count, requests, index, status)

- Parameters
  - count number of requests
  - requests array of requests
  - index index of request that completed
  - status status for the completed operations
- A call to MPI\_Waitany returns when one operation identified by the array of requests is complete

# Wait for non-blocking operations

`MPI_Waitsome(count, requests, done, index, status)`

- Parameters
  - `count` number of requests
  - `requests` array of requests
  - `done` number of completed requests
  - `index` array of indexes of completed requests
  - `status` array of statuses of completed requests
- A call to `MPI_Waitsome` returns when one or more operation identified by the array of requests is complete

# Non-blocking test for non-blocking operations

`MPI_Test(request, flag, status)`

- Parameters
  - `request` request
  - `flag` True if operation has completed status
  - `status` for the completed operations
- A call to `MPI_Test` is non-blocking. It allows one to schedule alternative activities while periodically checking for completion.

# Summary Non-blocking communication

- Non-blocking communication is usually the smarter way to do point-to-point communication in MPI
- Non-blocking communication realization
  - MPI\_Isend
  - MPI\_Irecv
  - MPI\_Wait(all)
- MPI-3 contains also non-blocking collectives (MPI\_Ibcast, MPI\_Ireduce, etc)

# Exercise: A dead-lock situation

- Safe Code:

```
if (my_rank==0)
{
MPI_Send( &tosend, 1, MPI_INT, 1, 10, MPI_COMM_WORLD);
MPI_Recv( &toreceive, 1, MPI_INT, 1, 20, MPI_COMM_WORLD, &status);
}
else if(my_rank == 1)
{
MPI_Recv( &toreceive, 1, MPI_INT, 0, 10, MPI_COMM_WORLD, &status);
MPI_Send( &tosend, 1, MPI_INT, 0, 20, MPI_COMM_WORLD);
}
```

# Exercise: A dead-lock situation

- Deadlock Code:

```
if (my_rank==0)
{
MPI_Recv( &toreceive, 1, MPI_INT, 1, 20, MPI_COMM_WORLD, &status);
MPI_Send( &tosend, 1, MPI_INT, 1, 10, MPI_COMM_WORLD);
}
else if(my_rank == 1)
{
MPI_Recv( &toreceive, 1, MPI_INT, 0, 10, MPI_COMM_WORLD, &status);
MPI_Send( &tosend, 1, MPI_INT, 0, 20, MPI_COMM_WORLD);
}
```

- Task:
  - Modify code send\_recv\_2b.c with non-blocking send and recv functions.
  - Don't forget to define "MPI\_Request send\_request,recv\_request;" first.

# Exercise: A dead-lock situation

- Buffering dependent Code

```
if (my_rank==0)
{
MPI_Send( &tosend, 1, MPI_INT, 1, 10, MPI_COMM_WORLD);
MPI_Recv( &toreceive, 1, MPI_INT, 1, 20, MPI_COMM_WORLD, &status);
}
else if(my_rank == 1)
{
MPI_Send( &tosend, 1, MPI_INT, 0, 20, MPI_COMM_WORLD);
MPI_Recv( &toreceive, 1, MPI_INT, 0, 10, MPI_COMM_WORLD, &status);
}
```

**Success of this code is dependent on buffering. One of the send must buffer and return. Otherwise, deadlock occurs.**



# Message Buffering

- Definition of “completion” for MPI\_Recv() is trivial – the data can now be used.
- Definition of “completion” for MPI\_Send() is trickier. Completion implies that the data has been stored away such that the program is free to overwrite the send “message” buffer.
- -- **Non-local**: the data can be sent directly to the receive buffer.
- -- **Local (buffering)**: the data can be stored in a local buffer (system provided or user provided), in which case the send could return before the receive is initiated.

# More about Communication Modes

Send Modes	MPI function	Completion Condition
Synchronous send	MPI_Ssend() MPI_Issend()	A send will not complete until a matching receive has been posted and the matching receive has begun reception of the data. Completion of a synchronous send not only indicates that the send buffer can be reused, but also indicates that the receiver has reached a certain point in its execution
Buffered send	MPI_Bsend() MPI_Ibsend()	Bsend() always completes (unless an error occurs) Completion is irrespective of the receiver.
Standard send	MPI_Send() MPI_Isend()	message sent (no guarantee that the receive has started). It is up to MPI to decide what to do.
Ready send	MPI_Rsend() MPI_Irsend()	may be used only when the a matching receive has already been posted

**“Recommendations:** In general, use MPI\_Send. If non-blocking routines are necessary, then try to use MPI\_Isend or MPI\_Irecv. Use MPI\_Bsend only when it is too inconvenient to use MPI\_Isend. The remaining routines, MPI\_Rsend, MPI\_Ssend, etc., are rarely used but may be of value in writing system-dependent message-passing code entirely within MPI.”

# Write Safe Code

- A safe MPI program should not rely on system buffering for success.
- Any system will eventually run out of buffer space as message sizes are increased.
- User should design proper send/receive orders to avoid **deadlock**.
- The user **must NOT** overwrite the send buffer until the send (data transfer) is complete.
- The user **can NOT** use the receiving buffer before the receive is complete.

# Exercise: Non-blocking 1d diffusion

```
initialize(uk, ukp1, numPoints, numProcs, myID);

for (step = 0; (step < maxsteps) && (maxdiff >= threshold); ++step) {

    double diff, maxdiff_local = 0.0;

    /* exchange boundary information */
    if (myID != 0)
        MPI_Send(&uk[1], 1, MPI_DOUBLE, leftNbr, 0, MPI_COMM_WORLD);
    if (myID != numProcs-1)
        MPI_Send(&uk[numPoints], 1, MPI_DOUBLE, rightNbr, 0, MPI_COMM_WORLD);
    if (myID != 0)
        MPI_Recv(&uk[0], 1, MPI_DOUBLE, leftNbr, 0, MPI_COMM_WORLD, &status);
    if (myID != numProcs-1)
        MPI_Recv(&uk[numPoints+1], 1, MPI_DOUBLE, rightNbr, 0, MPI_COMM_WORLD, &status);

    /* compute new values for interior points */
    for (i = 2; i < numPoints; ++i) {
        ukp1[i] = uk[i] + (dt/(dx*dx)) * (uk[i+1] - 2*uk[i] + uk[i-1]);
    }
    /* compute new values for boundary points
     * (no real need to do these separately, but it would allow us to
     * later overlap computation and communication with fewer code changes)
     */
    if (myID != 0) {
        int i = 1;
        ukp1[i] = uk[i] + (dt/(dx*dx)) * (uk[i+1] - 2*uk[i] + uk[i-1]);
    }
    if (myID != numProcs-1) {
        int i = numPoints;
        ukp1[i] = uk[i] + (dt/(dx*dx)) * (uk[i+1] - 2*uk[i] + uk[i-1]);
    }

    /* check for convergence */
}
```

# **Communicators and Groups**

# Communicators and Groups


- Many MPI users are only familiar with the communicator `MPI_COMM_WORLD`
- A **communicator** can be thought of a handle to a group
- A **group** is an ordered set of processes
- Each process is associated with a rank
- Ranks are contiguous and start from zero
- For many applications (dual level parallelism) maintaining different groups is appropriate
- Groups allow collective operations to work on a subset of processes
- Information can be added onto communicators to be passed into routines

# Communicators and Groups

- While we think of a communicator as spanning processes, it is actually unique to a process
- A communicator can be thought of as a handle to an object (group attribute) that describes a group of processes
- An **intracommunicator** is used for communication within a single group
- An **intercommunicator** is used for communication between 2 disjoint groups



# Communicators and Groups





# Communicators and Groups

- Refer to previous slide
- There are 4 distinct groups
- These are associated with **intracommunicators** MPI\_COMM\_WORLD, comm1, and comm2, and comm3
- $P_3$  is a member of 2 groups and may have different ranks in each group(say 3 & 4)
- If  $P_2$  wants to send a message to  $P_1$  it must use MPI\_COMM\_WORLD (intracommunicator) or comm5 (**intercommunicator**)
- If  $P_2$  wants to send a message to  $P_3$  it can use MPI\_COMM\_WORLD (send to rank 3) or comm1 (send to rank 4)
- $P_0$  can broadcast a message to all processes associated with comm2 by using **intercommunicator** comm4

# Example of using multiple communicators

Split a Large Communicator Into Smaller Communicators



```
MPI_Comm_split(  
    MPI_Comm comm,  
    int color,  
    int key,  
    MPI_Comm* newcomm)
```

- As the name implies, `MPI_Comm_split` creates new communicators by "splitting" a communicator into a group of sub-communicators based on the input values `color` and `key`.
- The first argument, `comm`, is the communicator that will be used as the basis for the new communicators. This could be `MPI_COMM_WORLD`, but it could be any other communicator as well.
- The second argument, `color`, determines to which new communicator each process will belong. All processes which pass in the same value for `color` are assigned to the same communicator.
- The third argument, `key`, determines the ordering (rank) within each new communicator. The process which passes in the smallest value for `color` will be rank 0, the next smallest will be rank 1, and so on.

# Example of using multiple communicators

```
int world_rank, world_size;
MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
MPI_Comm_size(MPI_COMM_WORLD, &world_size);

int color = world_rank / 4; // Determine color based on row

// Split the communicator based on the color and use the
// original rank for ordering
MPI_Comm row_comm;
MPI_Comm_split(MPI_COMM_WORLD, color, world_rank, &row_comm);

int row_rank, row_size;
MPI_Comm_rank(row_comm, &row_rank);
MPI_Comm_size(row_comm, &row_size);

printf("WORLD RANK/SIZE: %d/%d \t ROW RANK/SIZE: %d/%d\n",
       world_rank, world_size, row_rank, row_size);
```

# Other communicator creation functions

- `MPI_Comm_split` is the most common communicator creation function
- `MPI_Comm_dup` is the most basic and creates a duplicate of a communicator.
- `MPI_Comm_create` creates a new communicator with communication group
- A more flexible way to create communicators using a new kind of MPI object, `MPI_Group`

# Overview of groups

- **Groups vs. Communicators**

- A **group** is an ordered set of processes. Each process in a group is associated with a unique integer rank. A group is always associated with a **communicator** object.
- A **communicator** encompasses a **group** of processes that may communicate with each other. All MPI messages must specify a communicator. In the simplest sense, the communicator is an extra "tag" that must be included with MPI calls.
- From the programmer's perspective, **a group and a communicator are one**. The group routines are primarily used to specify which processes should be used to construct a communicator.

# Overview of groups

- **Groups** are dynamic - they can be created and destroyed during program execution.
- Processes may be in more than one group. They will have a unique rank within each group.
- MPI provides over **40** routines related to groups, communicators, and virtual topologies!
- Typical usage:
  1. Extract handle of global group from MPI\_COMM\_WORLD using MPI\_Comm\_group
  2. Form new group as a subset of global group using **MPI\_Group\_incl**
  3. Create new communicator for new group using **MPI\_Comm\_create**
  4. Determine new rank in new communicator using **MPI\_Comm\_rank**
  5. Conduct communications using any MPI message passing routine
  6. When finished, free up new communicator and group (optional) using **MPI\_Comm\_free** and **MPI\_Group\_free**

# Example of using groups

```
// Get the rank and size in the original communicator
int world_rank, world_size;
MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
MPI_Comm_size(MPI_COMM_WORLD, &world_size);

// Get the group of processes in MPI_COMM_WORLD
MPI_Group world_group;
MPI_Comm_group(MPI_COMM_WORLD, &world_group);

int n = 7;
const int ranks[7] = {1, 2, 3, 5, 7, 11, 13};

// Construct a group containing all of the prime ranks in world_group
MPI_Group prime_group;
MPI_Group_incl(world_group, 7, ranks, &prime_group);

// Create a new communicator based on the group
MPI_Comm prime_comm;
MPI_Comm_create_group(MPI_COMM_WORLD, prime_group, 0, &prime_comm);

int prime_rank = -1, prime_size = -1;
// If this rank isn't in the new communicator, it will be
// MPI_COMM_NULL. Using MPI_COMM_NULL for MPI_Comm_rank or
// MPI_Comm_size is erroneous
if (MPI_COMM_NULL != prime_comm) {
    MPI_Comm_rank(prime_comm, &prime_rank);
    MPI_Comm_size(prime_comm, &prime_size);
}


printf("WORLD RANK/SIZE: %d/%d \t PRIME RANK/SIZE: %d/%d\n",
       world_rank, world_size, prime_rank, prime_size);
```



# Using MPI groups


Once you have a group or two, performing operations on them is straightforward. Getting the union looks like this:

```
MPI_Group_union(  
    MPI_Group group1,  
    MPI_Group group2,  
    MPI_Group* newgroup)
```



And you can probably guess that the intersection looks like this:

```
MPI_Group_intersection(  
    MPI_Group group1,  
    MPI_Group group2,  
    MPI_Group* newgroup)
```



In both cases, the operation is performed on `group1` and `group2` and the result is stored in `newgroup`.

# Virtual and Physical Topologies

- A *virtual topology* represents the way that MPI processes communicate
  - Nearest neighbor exchange in a mesh
  - recursive doubling in an all-to-all exchange
- A *physical topology* represents that connections between the cores, chips, and nodes in the hardware

# MPI's Topology Routines

- MPI provides routines to create new **intra-communicators** that order the process ranks in a way that *may* be a better match for the *physical topology*
- Two types of virtual topology supported:
  - Cartesian (regular mesh)
  - Graph (several ways to define in MPI)
- Additional routines provide access to the defined virtual topology
- (Virtual) topologies are properties of a communicator
  - Topology routines all create a *new* communicator with properties of the specified virtual topology

# MPI Cartesian Topology

- Example: 12 processes arranged on a 3 x 4 grid

0 (0,0)	1 (0,1)	2 (0,2)	3 (0,3)
4 (1,0)	5 (1,1)	6 (1,2)	7 (1,3)
8 (2,0)	9 (2,1)	10 (2,2)	11 (2,3)

# MPI Cartesian Topology

- Create a new virtual topology using
  - MPI\_Cart\_create
- Determine “good” sizes of mesh with
  - MPI\_Dims\_create

# MPI Cartesian Topology

```
int MPI_Cart_create (MPI_Comm comm_old , int ndims ,  
int *dims , int *periods , int reorder , MPI_Comm  
*comm_cart)
```

- Creates a new communicator `comm_cart` from `comm_old`, that represents an `ndims` dimensional mesh with sizes `dims`.
- The mesh is periodic in coordinate direction if `periods[i]` is true.
- The ranks in the new communicator are reordered (to better match the physical topology) if `reorder` is true (Set to false, rank in `comm_cart` must be the same as in `comm_old`)

# MPI Cartesian Topology

- Creates logical 2-d Mesh of size 3x4

```
int dims[2] = {3,4};
int periods[2] = {0,0};
MPI_Comm topocomm;
MPI_Cart_create(comm, 3, dims, periods, 0,
&topocomm);
```
- But we're starting MPI processes with a one-dimensional argument (-p X)
  - User has to determine size of each dimension
  - Often as “square” as possible, MPI can help!

# MPI Cartesian Topology

- `MPI_Dims_create(int nnodes, int ndims, int *dims)`
  - Fill in the `dims` array such that the product of `dims[i]` for  $i=0$  to  $ndims-1$  equals `nnodes`.
  - Any value of `dims[i]` that is 0 on input will be replaced; Non-zero entries in `dims` will not be changed
  - `MPI_Dims_create(12, 2, dims) == (3,4) or (4,3)?`



# CARTESIAN QUERY FUNCTIONS

- `MPI_Cartdim_get()`
  - Gets dimensions of a Cartesian communicator
- `MPI_Cart_get()`
  - Gets size of dimensions
- `MPI_Cart_rank()`
  - Translate coordinates to rank
- `MPI_Cart_coords()`
  - Translate rank to coordinates

# CARTESIAN COMMUNICATION HELPERS

```
MPI_Cart_shift(MPI_Comm comm, int  
direction, int disp,  
int *rank_source, int *rank_dest)
```

- Returns the ranks of the processes that are a shift of **disp** steps in coordinate **direction**
- Useful for nearest neighbor communication in the coordinate directions
- May return MPI\_PROC\_NULL

# MPI Graph Topology

- MPI provides routines to specify a general graph virtual topology
  - Graph vertices represent MPI processes (usually one per process)
  - Graph edges indicate important connections (e.g., nontrivial communication between the connected processes)
  - Edge weights provide more information (e.g., amount of communication)

# MPI Graph Topology

```
int MPI_Graph_create(MPI_Comm comm_old, int  
nnodes, int *index, int *edges, int reorder,  
MPI_Comm *comm_graph)
```

- nnodes is the total number of nodes
- index i stores the total number of neighbors for the first i nodes (sum)
  - Acts as offset into edges array
- edges stores the edge list for all processes
  - Edge list for process j starts at index[j] in edges
  - Process j has index[j+1]-index[j] edges

# DISTRIBUTED GRAPH CONSTRUCTOR

- `MPI_Graph_create` is discouraged
  - Not scalable
  - Not deprecated yet but hopefully soon
- New distributed interface:
  - Scalable, allows distributed graph specification
    - Either local neighbors or any edge in the graph
  - Specify edge weights
    - Meaning undefined but optimization opportunity for vendors!
- Info arguments
  - Communicate assertions of semantics to the MPI library
  - E.g., semantics of edge weights

# DISTRIBUTED GRAPH CONSTRUCTOR

`MPI_Dist_graph_create_adjacent(MPI_Comm oldcomm, int indegree, int sources[], int sourceweights[], int outdegree, int dests[], int destweights[], MPI_Info info, int qreorder, MPI_Comm *newcomm)`

- Describe *only* the graph vertex corresponding to the calling process
  - Hence “Dist\_graph” – distributed description of graph
- Graph is directed – separate in and out edges
- **info** allows additional, implementation-specific information
- **qreorder** if true lets MPI implementation reorder ranks for a better mapping to physical topology
- **MPI\_UNWEIGHTED** may be used for weights *arrays*

# Other Graph Routines

- `MPI_Dist_graph_create`
  - More general, allows multiple graph vertices per process
- Information on graph
  - `MPI_Dist_graph_neighbors_count`
    - Query the number of neighbors of calling process
  - `MPI_Dist_graph_neighbors`
    - Query the neighbor list of calling process
- Collective Communication along arbitrary neighborhoods
  - `MPI_NEIGHBOR_ALLGATHER`
  - `MPI_NEIGHBOR_ALLTOALL`

# Topology: good and bad

- A common virtual topology is *nearest neighbor in a mesh*
  - Matrix computations
  - PDE Simulations on regular computational grids
- Many Large Scale Systems use a mesh as the physical topology
  - IBM Blue Gene series; Cray through XE6/XK7
- Performance can depend on how well the virtual topology is mapped onto the physical topology
- Bisection bandwidth does *not* scale with network size of a mesh/torus network
- Non-nearest neighbor communication suffers from contention
- Graph topology: Complex to implement. No good implementations in general use; research work limited



# **ONE-SIDED COMMUNICATION**

# ONE-SIDED COMMUNICATION

- Two components of message-passing: sending and receiving
- One-sided communication
  - Only single process calls data movement functions (put or get) – remote memory access (RMA)
  - Communication patterns specified by only a single process – Always non-blocking

# Why one-sided communication?

- Certain algorithms featuring unstrucutred communication easier to implement
- Potentially reduced overhead and improved scalability
- Hardware support for remote memory access has been restored in most current-generation architectures

# Origin and target

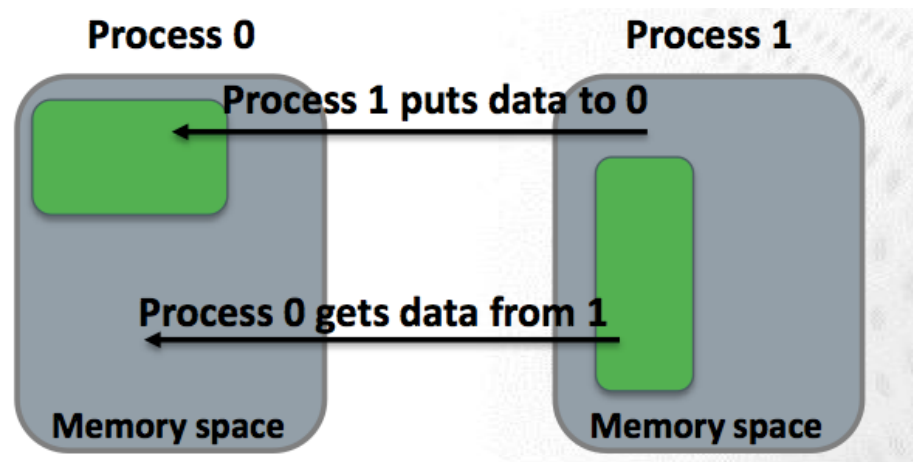
- Origin process: a process which calls data movement function
- Target process: a process whose memory is accessed

# Origin and target

- Origin process: a process which calls data movement function
- Target process: a process whose memory is accessed

# Remote memory access window

- **Window** is a region in process's memory which is made available for remote operations
- Windows are created by collective calls
- Windows may be different in different processes



# Data movement operations in MPI

- PUT data to the memory in target process
  - From local buffer in origin to the window in target
- GET data from the memory of target process
  - From the window in target to the local buffer in origin
- ACCUMULATE data in target process
  - Use local buffer in origin and update the data (e.g. add the data from origin) in the window in target
  - One-sided reduction

# Synchronization

- Communication takes place within epochs
  - Synchronization calls start and end an epoch
  - There can be multiple data movement calls within an epoch
  - An epoch is specific to a particular window
- Active synchronization:
  - Both origin and target perform synchronization calls
- Passive synchronization:
  - No MPI calls at target process



# One-sided communication in a nutshell

- Define a memory window Start an epoch
  - Target: exposure epoch
  - Origin: access epoch
- GET, PUT, and/or ACCUMULATE data
- Complete the communications by ending the epoch

# Creating a window

`MPI_Win_create(base, size, disp_unit, info, comm, win)`

- base (pointer to) local memory to expose for RMA
- size size of a window in bytes
- disp\_unit local unit size for displacements in bytes
- info hints for implementation
- comm communicator
- win handle to window

# Starting and ending an epoch

`MPI_Win_fence(assert, win)`

- assert optimize for specific usage. Valid values are "0", `MPI_MODE_NOSTORE`, `MPI_MODE_NOPUT`, `MPI_MODE_NOPRECEDE`, `MPI_MODE_NOSUCCEED`
- win window handle
- Used both for starting and ending an epoch
  - Should both precede and follow data movement calls
- Collective, barrier-like operation

# Data movement: Put

`MPI_Put(origin, origin_count, origin_datatype, target_rank, target_disp, target_count, target_datatype, win)`

- `origin` (pointer to) local data to be send to target
- `origin_count` number of elements to put
- `origin_datatype` MPI datatype for local data
- `target_rank` rank of the target task
- `target_disp` starting point in target window
- `target_count` number of elements in target
- `target_datatype` MPI datatype for remote data
- `win` RMA window

# Simple example: Put

```
...
int data;
MPI_Win window;
data = rank;

// Create window
MPI_Win_create(&data, sizeof(int), sizeof(int), MPI_INFO_NULL, MPI_COMM_WORLD,
&window);

...
MPI_Win_fence(0, window);

if (rank == 0)
    MPI_Put(&data, 1, MPI_INT, 1, 0, 1, MPI_INT, window);

MPI_Win_fence(0, window);
...
MPI_Win_free(&window);
```

# Performance considerations/summary

- Performance of the one-sided approach is highly implementation-dependent
- Maximize the amount of operations within an epoch
- Provide the assert parameters for MPI\_Win\_fence
- One-sided communication allows communication patterns to be specified from a single process
- Can reduce synchronization overheads and provide better performance especially on recent hardware
- Basic concepts:
  - Creation of the memory window
  - Communication epoch
  - Data movement operations (MPI\_Put, MPI\_Get etc)